An Intelligent Energy Brain for Smart Buildings

Instructor : Br. Bhaswarachaitanya (Tamal Maharaj)

Jahar Kumar Paul (B2430049)

Ayan Kumar Batabyal (B2530074)

Sayan Goswami (B2530098)

Introduction

Motivation:

"We do not inherit the Earth from our ancestors; we borrow it from our children." — American Proverb

Sustainable living begins with intelligent choices. As the world's energy demand continues to rise, it becomes essential to design systems that not only deliver comfort but also preserve the environment. Smart buildings, powered by intelligent systems, offer a path toward this greener future.

Problem Description

We aim to design an **intelligent energy management system** that combines **Machine Learning (ML)** and **Reinforcement Learning (RL)** to make buildings truly smart — ensuring comfort while optimizing energy consumption.

Air conditioning systems are among the highest energy consumers in commercial spaces. Our focus is on **reducing HVAC energy usage** without compromising comfort.

The proposed **Hybrid Model** functions in two major phases:

1. **Deep Learning Phase** – for prediction and forecasting.
2. **Reinforcement Learning Phase** – for adaptive control and real-time decision-making.

Proposed Approach and Algorithms

Consider a scenario where a client meeting is scheduled at **11:00 AM** in a large corporate office that starts operations at **10:00 AM**. To ensure maximum comfort for visiting clients, the air conditioning system must start beforehand.

However, determining the appropriate temperature setting depends on several factors:

- Room size

- Number of occupants
- Outdoor temperature
- Current indoor temperature
- Energy price
- Desired comfort level

And most importantly — energy must be used efficiently.

Phase 1: Deep Learning Prediction

We begin by training a **Multi-Output Regressor(Deep Neural Network)** using historical building data.

Input Features: Room size, occupancy, outdoor temperature, current indoor temperature, energy price, etc.

Model Outputs: Predicted values for:

- Indoor temperature
- Outdoor temperature
- Forecasted occupancy

These predictions are generated for **11:30 AM**, providing the RL model with a short-term forecast window.

Phase 2: Reinforcement Learning Control (11:00 AM – 11:30 AM)

The **Reinforcement Learning Model** continuously observes the environment — including sensor data, predictions from the neural network, and human feedback.

For example:

- If the forecasted occupancy changes (more or fewer people than expected), the RL model dynamically adjusts the temperature to maintain comfort while minimizing energy use.
- If someone manually changes the temperature using the remote:
  - A **large penalty** is applied for temperature changes (indicating discomfort).
  - A **smaller penalty** is applied for fan speed adjustments.

- If no manual change occurs, the system receives a **reward**, implying satisfactory comfort levels.

- The RL model also monitors environmental feedback, such as people entering or leaving the room or sudden weather changes (e.g., unexpected rain).

This data is collected every **7–8 minutes**, allowing the model to adapt in near real time.

Dynamic Update (11:30 AM – 11:35 AM)

If there are changes in occupancy, indoor temperature, or outdoor temperature during this window, the **Deep Learning model's weights** are updated. This ensures that the system learns continuously and remains aligned with real-world conditions.

Loophole and Tackling it

There exists a potential gap in the proposed framework till now.

Who sets the initial AC temperature — manually or automatically?

Both approaches are feasible. However, to maintain system autonomy and consistency, we propose employing a **lightweight, pre-trained Reinforcement Learning (RL) model** to determine the initial temperature setting.

This initial RL model will act as a warm-start mechanism, selecting a comfortable yet energy-efficient temperature before the main control phase begins. Over time, as more feedback is collected, the model can be fine-tuned to improve its initial decision-making accuracy.

Outcome

By integrating **Deep Learning for prediction** and **Reinforcement Learning for adaptive control**, the system will:

- Significantly reduce HVAC and overall energy consumption.
- Maintain optimal indoor comfort.
- Continuously learn from feedback and environmental changes.
- Operate autonomously and scale efficiently.

Ultimately, this intelligent framework will act as an **energy brain** for smart buildings — making data-driven decisions that balance **comfort, cost, and sustainability**.

Tentative Work Division

Machine and Deep Learning - JAHAR KUMAR PAUL
Reinforcement Learning - AYAN KUMAR BATABYAL & SAYAN GOSWAMI